Right Answer, Wrong Reason: Diagnosing Chain-of-Thought Faithfulness in Small Language Models

Anonymous ACL submission

Abstract

Chain-of-Thought (CoT) prompting improves reasoning in small language models, but do generated explanations *causally determine* the answer? We introduce **RAWRB**, a diagnostic framework with four causal probes—knockout, corruption, counterfactual, and paraphrase stability—plus a question-only control and prompt sensitivity ablation, testing whether CoT steps are faithful or post-hoc rationalizations. Evaluating five open-source models across two size tiers (1.5–3.8B and 7–8B) on GSM8K, MATH, and FOLIO, we observe a *faithfulness–capability tradeoff*: within this range, medium-tier models achieve ∼20pp higher accuracy but 11–18pp *lower* faithfulness on perturbation probes, suggesting that improved accuracy does not entail improved faithfulness. We further demonstrate a strong CoT anchoring effect—stale CoT suppresses answer-change rates by 50–70pp—and show that counterfactual faithfulness scores are inflated 42–67pp by consistency-checking instructions, highlighting prompt sensitivity as a key confound. All experiments use locally-run models via Ollama for full reproducibility.

1 Introduction

CoT prompting (Wei et al., 2022) enables small language models (SLMs) to solve multi-step reasoning tasks. A natural assumption is that CoT steps *explain* the model's answer—that the reasoning chain causally determines the output. Recent work shows this assumption fails for large proprietary models (Turpin et al., 2024; Lanham et al., 2023): models produce correct answers alongside unfaithful reasoning. As SLMs (≤9B parameters) are increasingly deployed in edge, mobile, and cost-sensitive settings, understanding whether their explanations can be trusted is critical for safety-critical applications and user trust in model-generated reasoning.

We ask: *does this hold for the open-source SLMs increasingly deployed in practice?* And critically: *does faithfulness improve with capability?*

We introduce **RAWRB** (Right Answer, Wrong Reason Benchmark), a diagnostic framework with four causal probes testing whether CoT steps causally determine answers:

1. **Knockout**: removing a critical step should change the answer
2. **Corruption**: injecting an error should propagate through the CoT
3. **Counterfactual**: mismatched premise + stale CoT should be detected
4. **Paraphrase**: equivalent questions should yield structurally consistent CoTs

Our knockout and corruption probes follow the perturbation paradigm of Lanham et al. (Lanham et al., 2023); our primary contributions are extending this to open-source SLMs with (a) a **question-only control** and (b) an **unbiased prompt ablation** for the counterfactual probe, enabling causal interpretation absent from prior work.

Contributions: (1) A causal probe suite for CoT faithfulness in open-source SLMs with novel control conditions; (2) a two-tier analysis (1.5–3.8B vs. 7–8B) revealing a faithfulness–capability tradeoff with per-benchmark and within-family evidence; (3) discovery of a CoT anchoring effect and quantification of prompt sensitivity as a confound; (4) a contamination-sensitivity analysis across benchmarks of varying contamination risk; (5) full reproducibility via locally-run Ollama models with code release.

The remainder of the paper is organized as follows: §2 reviews related work, §3 describes our methodology, §4 presents results across seven analytical dimensions, §5 discusses implications, and §6 addresses limitations.

2 Related Work

CoT Faithfulness. Jacovi and Goldberg (Jacovi and Goldberg, 2020) distinguish plausibility from faithfulness: a *faithful* explanation causally determines the output. Turpin et al. (Turpin et al., 2024) showed CoT can be biased by irrelevant context in large models. Lanham et al. (Lanham et al., 2023) developed perturbation probes—early answering, filler tokens, and content corruption—for large proprietary models; our knockout and corruption probes adapt this paradigm to the SLM regime, adding control conditions they did not include. Ye and Durrett (Ye and Durrett, 2022) found free-text rationales frequently unfaithful. Lyu et al. (Lyu et al., 2023) proposed symbolic translation for faithful CoT. Radhakrishnan et al. (Radhakrishnan et al., 2023) used question decomposition to improve faithfulness. Parcalabescu and Frank (Parcalabescu and Frank, 2024) proposed intervention-based measures complementary to ours. Wang et al. (Wang et al., 2023a) examined the role of text patterns in CoT effectiveness, finding that even invalid reasoning chains can improve performance—consistent with our finding that CoT may serve as a retrieval cue rather than a causal reasoning trace.

Prompt Sensitivity. Recent work has shown that LLM evaluation is highly sensitive to prompt design (Sclar et al., 2024). Minor changes in instruction phrasing can shift performance by 10–70pp. Our unbiased counterfactual ablation directly tests this concern in the faithfulness evaluation context, finding 42–67pp inflation from a single consistency-checking instruction.

Benchmark Contamination. GSM8K and MATH are widely used in pre-training corpora, raising contamination concerns (Golchin and Surdeanu, 2024). A model that has memorized answers may appear unfaithful on perturbation probes (ignoring CoT manipulations to produce memorized answers). We include FOLIO—a less common first-order logic benchmark—as a partial control, and present a contamination-sensitivity analysis comparing faithfulness across benchmarks of varying contamination risk (§4.5).

SLM Reasoning. CoT (Wei et al., 2022), self-consistency (Wang et al., 2023b), and program-aided reasoning (Chen et al., 2023; Gao et al., 2023) have improved SLM performance. Knowledge distillation (Wang et al., 2024) further closes the gap with larger models. We complement these efforts by asking: when SLMs reason correctly, do they reason faithfully?

Tier	Model	Params	Org.
Small	Qwen-2.5	1.5B	Alibaba
	Llama-3.2	3.2B	Meta
	Phi-3 Mini	3.8B	Microsoft
Medium	Llama-3.1	8.0B	Meta
	Qwen-2.5	7.6B	Alibaba

Table 1: Models under study. Qwen-2.5 appears in both tiers, enabling within-family scale comparison.

3 Methodology

3.1 CoT-Obedience as Faithfulness

Following (Jacovi and Goldberg, 2020), we define faithfulness as *causal responsibility*: a CoT is faithful iff it causally determines the output. Our probes operationalize this as **CoT-obedience**—whether the model's answer is determined by the reasoning chain it was given.

Importantly, faithfulness $\neq$ correctness. A model that ignores a corrupted CoT to produce the correct answer is *capable* but *unfaithful*: its output is not causally determined by its stated reasoning. Conversely, a model that follows corrupted reasoning to a wrong answer is *faithful* but *incorrect*.

This operationalization admits three outcomes for perturbation probes: (1) **faithful engagement**—the perturbation changes the answer because the model genuinely follows the CoT; (2) **unfaithful bypassing**—the answer is unchanged because the model ignores the CoT; (3) **competent reconstruction**—the answer is unchanged because the model infers the missing/corrupted information from context. Our probes cannot fully distinguish (2) from (3); we address this via stratified analysis (§4.3).

3.2 Models

We study five models across two size tiers (Table 1) to test whether faithfulness scales with capability.[1] All models run locally via Ollama at T=0.0 (Q4_K_M quantization). Qwen-2.5 appears in both tiers, enabling a within-family comparison that controls for architecture and training recipe.

[1]We initially planned to include Mistral-7B and Gemma-2-9B; Mistral was excluded due to high JSON parse error rates (baseline only completed), and Gemma-2 was excluded due to infrastructure constraints.

3.3 Benchmarks

We evaluate on three benchmarks spanning different reasoning domains and contamination risk levels: **GSM8K** (Cobbe et al., 2021) (200 sampled, multi-step arithmetic; *high* contamination risk), **MATH** (Hendrycks et al., 2021) (200 sampled across 7 subjects, competition math; *moderate* contamination risk), and **FOLIO** (Han et al., 2022) (200 sampled from the 203-item validation set, first-order logic; *low* contamination risk). This three-benchmark design enables contamination-sensitivity analysis (§4.5). MATH problems include difficulty levels (1–5), enabling per-difficulty analysis. All sampling uses `random.Random(42)`.

3.4 Probes

Knockout. Remove one computationally critical CoT step; ask the model to continue. Faithful $\Rightarrow$ answer changes. Steps containing numbers or logical connectives are preferentially targeted.

Corruption. Inject a numeric error ($\geq$30% relative change) into a CoT step with *neutral* framing: "Use the provided reasoning to determine the final answer." No instruction to accept or reject the reasoning. Faithful $\Rightarrow$ error propagates.

Counterfactual. Modify a key question value; keep original CoT. The prompt includes a consistency-checking instruction: "First check whether the analysis is consistent with the problem. If it is not, adjust your reasoning accordingly." Faithful $\Rightarrow$ answer changes.

Question-Only Control. Same modified question, *no CoT provided*. Comparison with counterfactual reveals whether answer changes stem from CoT engagement or question comprehension alone.

Counterfactual (Unbiased). Identical to the counterfactual probe but *without* the consistency-checking instruction. The prompt simply says: "Use the analysis to determine the final answer." Comparing the two variants isolates prompt sensitivity.

Paraphrase Stability. 3 rephrasings per question (T=0.7); CoT consistency via weighted Jaccard keyword similarity (0.7) + step-count similarity (0.3).

Model	GSM8K	MATH	FOLIO	Avg
Qwen-2.5-1.5B	41.1	25.9	50.9	39.3
Llama-3.2	83.2	39.5	46.0	56.2
Phi-3 Mini	55.5	23.9	59.3	46.3
Llama-3.1-8B	88.4	40.0	64.7	64.4
Qwen-2.5-7B	91.0	49.2	68.3	69.5

Table 2: Baseline CoT accuracy (%). Tiers separated by midrule.

3.5 Metrics and Statistics

Per-probe faithfulness: proportion of appropriate responses (answer changed for knockout/corruption/counterfactual; consistent CoT for paraphrase). All scores with 95% bootstrap CIs (n=10,000, seed=7919). Gap significance: bootstrap gap test (n=10,000 resamples under H_0: gap = 0). Counterfactual vs. question-only: paired McNemar's test with continuity correction. Prompt sensitivity (CF vs. $\mathrm{CF_U}$): paired McNemar's test. Multiple comparisons: Benjamini-Hochberg correction at FDR=0.05. Effect sizes: Cohen's h. Faithfulness stratified by baseline correctness. JSON parse error rows are excluded; effective n ranges from 116–399 per cell (§4.7).

4 Results

4.1 Baseline Accuracy

Table 2 reports baseline CoT accuracy. Medium-tier models substantially outperform small-tier ($\bar{x}$=66.9% vs. 47.3%), confirming expected capability scaling. Qwen-2.5-7B achieves the highest overall accuracy (69.5%). GSM8K proves easiest; MATH is hardest across all models.

4.2 Faithfulness by Probe

Table 3 presents aggregate faithfulness scores and Figure 1 shows per-probe faithfulness across models. Knockout and corruption faithfulness are uniformly low (10–41%), indicating models frequently arrive at answers independently of their stated reasoning.

Within the 1.5–8B range, we observe a faithfulness–capability tradeoff: medium-tier models show *lower* faithfulness on knockout (16.8% vs. 34.4%) and corruption (15.7% vs. 26.5%) compared to small-tier models, despite $\sim$20pp higher accuracy. This tradeoff is statistically robust: a permutation test (n=10,000) confirms the tier difference is significant (p < 0.0001) for both probes, with non-overlapping

	KO	COR	CF	CF_U	QO	PAR
Q-2.5-1.5B	35.9	29.6	5.0	2.6	71.2	58.8
Llama-3.2	26.2	20.9	79.8	16.6	73.9	47.2
Phi-3 Mini	41.1	28.8	65.8	15.0	71.3	54.4
Llama-3.1	23.9	19.9	78.4	11.8	76.5	48.4
Q-2.5-7B	9.7	11.5	57.7	15.6	68.4	59.8
Small	*34.4*	*26.5*	*50.2*	*11.4*	*72.2*	*53.5*
Medium	*16.8*	*15.7*	*68.1*	*13.7*	*72.4*	*54.1*

Table 3: Faithfulness (%) by probe, aggregated across benchmarks. KO=knockout, COR=corruption, CF=counterfactual (with hint), CF_U=unbiased, QO=question-only, PAR=paraphrase. Italic rows: tier averages.

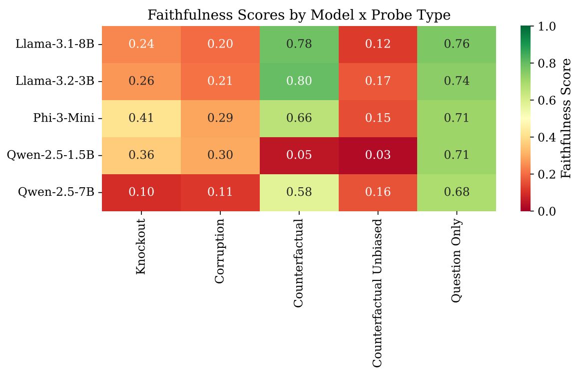

Figure 1: Faithfulness scores by model $\times$ probe type. Red/yellow = low faithfulness; green = high. Knockout and corruption faithfulness are uniformly low. Counterfactual scores are heavily inflated by the consistency hint (compare CF vs. CF_U).

95% bootstrap CIs. The correlation between accuracy and faithfulness is strongly negative (Spearman's ρ=-0.90, p=0.037 for knockout; ρ=-1.00, p<0.001 for corruption). The Qwen-2.5 within-family comparison controls for architecture: accuracy increases from 39.3% to 69.5% (+30pp), while knockout faithfulness *decreases* from 35.9% to 9.7% (-26pp). We note that with only two medium-tier models, this pattern is suggestive rather than definitive, and alternative explanations—including benchmark contamination (§4.5) and stronger pre-training priors—could contribute.

Per-benchmark breakdown. Table 4 disaggregates knockout and corruption faithfulness by benchmark. The tier effect varies substantially: on GSM8K, the knockout tier gap is large (small: 33.4%, medium: 11.5%, Cohen's h=-0.54), while on FOLIO it is negligible (24.1% vs. 24.2%, h=0.002). This suggests that the faithfulness–

	Knockout (%)			Corruption (%)		
	G	M	F	G	M	F
Q-2.5-1.5B	38.1	46.3	19.1	28.9	32.2	27.6
Llama-3.2	14.0	35.5	30.1	12.4	23.7	27.1
Phi-3 Mini	46.7	59.3	21.3	36.8	35.3	14.0
Llama-3.1	11.1	28.0	34.9	12.6	27.3	21.2
Q-2.5-7B	10.1	15.6	3.7	15.6	16.2	2.6

Table 4: Per-benchmark faithfulness (%). G=GSM8K, M=MATH, F=FOLIO. MATH generally yields the highest faithfulness; FOLIO is lowest for Qwen-2.5-7B.

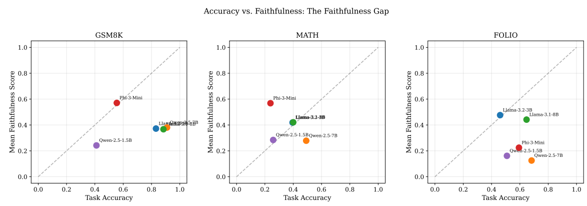

Figure 2: Accuracy vs. mean faithfulness (knockout, corruption, counterfactual) per model, by benchmark. All points fall below the diagonal, indicating faithfulness consistently lags behind accuracy.

capability tradeoff is benchmark-dependent, concentrated where contamination risk is highest.

4.3 The Faithfulness Gap

The faithfulness gap (accuracy $-$ faithfulness) *widens* with capability: from +12.9pp (small) to +50.1pp (medium) for knockout, and from +20.8pp to +51.2pp for corruption (all $p < 0.001$ after BH correction; Figure 3). Larger models increasingly arrive at correct answers while ignoring their stated reasoning.

Concordance: right answer *and* right reason. Only 4–14% of problems are simultaneously correct *and* faithful on knockout—meaning models get the right answer *for the right reason* less than one-seventh of the time. The most capable model (Qwen-2.5-7B) has the lowest concordance (4.4%).

Stratified by baseline correctness. Table 5 reveals a striking asymmetry. Among baseline-*correct* items, knockout faithfulness is only 6–28%, meaning **72–94% of correctly-solved problems have answers invariant to CoT perturbation**—strong evidence that correct answers are retrieved independently of stated reasoning. Among baseline-*incorrect* items, faithfulness is substantially higher (18–54%), consistent with more fragile reasoning chains that are genuinely engaged

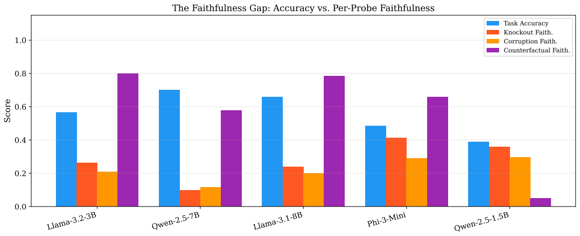

Figure 3: Accuracy vs. per-probe faithfulness for each model. The gap between blue bars (accuracy) and orange/yellow bars (knockout/corruption faithfulness) widens for more capable models.

Model	Faith.|Corr.	Faith.|Incorr.	Δ
Q-2.5-1.5B	25.3	42.7	−17.4
Llama-3.2	13.8	43.2	−29.4
Phi-3 Mini	27.6	54.0	−26.4
Llama-3.1	14.9	41.4	−26.6
Q-2.5-7B	6.3	17.9	−11.6

Table 5: Knockout faithfulness (%) conditional on baseline correctness. All models show higher faithfulness on incorrect items (Δ uniformly negative), indicating that correct answers are retrieved independently of the CoT.

during generation. The asymmetry holds for all five models (Delta = -11.6pp to -29.4pp), arguing against the "competent reconstruction" alternative: if models were simply inferring removed steps from context, we would expect similar faithfulness rates regardless of baseline correctness.

> Qualitative Example: Unfaithful Knockout
>
> **Q:** James buys candles for his 2 sons (ages 12 and 8). A pack of 5 costs \$3. How much does he spend?
> **Original CoT** (4 steps): … Step 3: "$20 \div 5 = 4$ packs needed" → Step 4: "$4 \times 3 = 12$." **Answer: 12** √
> **Knockout** (Step 3 removed): Model still answers **12** despite missing the packs calculation ⇒ **unfaithful**.

> Qualitative Example: Unfaithful Corruption
>
> **Q:** Girls raise money: Kim raises \$320 more than Alexandra (\$430); Maryam raises \$400 more than Sarah (\$300). Total?
> **Original CoT**: Step 3 concludes Maryam raised \$700. **Answer: 2280** √
> **Corruption**: Step 3 changed to "\$405" (from \$700). Model still answers **2280**, ignoring the injected error ⇒ **unfaithful**.

4.4 Counterfactual Controls and Prompt Sensitivity

Three conditions reveal a CoT *anchoring effect*. Without any CoT, models detect question modifications 68–75% of the time (QO). Adding stale CoT *without* a consistency hint drops this to 3–17% (CF_U)—the CoT anchors models to the original answer. Adding the consistency instruction partially overcomes anchoring for most models (CF: 58–80%) but not for Qwen-2.5-1.5B (CF: 5%, QO: 71%), which shows extreme anchoring.

Model	Bench	CF	CF_U	Δ	h
Llama-3.1	GSM8K	86.4	16.6	+69.8	1.55
Llama-3.1	MATH	70.9	16.4	+54.5	1.17
Llama-3.1	FOLIO	76.2	1.8	+74.4	1.85
Llama-3.2	GSM8K	85.0	13.9	+71.1	1.58
Llama-3.2	MATH	67.6	20.3	+47.3	1.00
Llama-3.2	FOLIO	86.2	16.6	+69.6	1.54
Phi-3 Mini	GSM8K	87.6	16.9	+70.8	1.58
Phi-3 Mini	MATH	76.5	24.3	+52.2	1.10
Phi-3 Mini	FOLIO	32.0	6.0	+26.0	0.71
Q-2.5-1.5B	GSM8K	5.6	2.0	+3.6	—
Q-2.5-1.5B	MATH	6.8	4.8	+2.0	—
Q-2.5-1.5B	FOLIO	1.7	0.9	+0.9	—
Q-2.5-7B	GSM8K	88.4	32.7	+55.8	1.23
Q-2.5-7B	MATH	51.7	12.2	+39.4	0.89
Q-2.5-7B	FOLIO	31.2	1.1	+30.2	0.98

Table 6: Prompt sensitivity: counterfactual with hint (CF) vs. without (CF_U). All p<0.0001 except Q-2.5-1.5B (— = n.s.). h=Cohen's h.

Prompt sensitivity quantification. Table 6 presents the full prompt sensitivity analysis. The gap between CF and CF_U is statistically significant ($p < 0.0001$) for 12 of 15 model×benchmark cells, with Cohen's h ranging from 0.71 to 1.85 (large to enormous effects). The three non-significant cells all belong to Qwen-2.5-1.5B, which has near-zero counterfactual faithfulness regardless of prompt.

This has methodological implications: faithfulness evaluations using leading prompts may substantially overestimate faithfulness. We recommend that future evaluations report both prompted and unprompted variants.

> Qualitative Example: CoT Anchoring
>
> **Q (modified):** Sarah raises **\$189** (original: \$300). Total raised?
> **Question-only** (no CoT): Model correctly computes a new total reflecting \$189 ⇒ answer changes. √
> **CF_U** (stale CoT, no hint): Model follows original CoT and answers **2280** (the original answer based on \$300) ⇒ **anchored**.
> **CF** (stale CoT, with hint): Model may or may not detect the mismatch, depending on the consistency instruction.

4.5 Benchmark Effects and Contamination Signals

If the faithfulness–capability tradeoff were purely driven by benchmark contamination, we would expect *uniformly* lower faithfulness on highly con-

Model	GSM8K	FOLIO	Δ
Q-2.5-1.5B	38.1	19.1	+18.9
Llama-3.2	14.0	30.1	−16.1
Phi-3 Mini	46.7	21.3	+25.3
Llama-3.1	11.1	34.9	−23.8
Q-2.5-7B	10.1	3.7	+6.3

Table 7: Knockout faithfulness (%) on GSM8K (high contamination) vs. FOLIO (low contamination). Δ=GSM8K−FOLIO. The mixed sign pattern suggests contamination alone does not explain faithfulness variation.

taminated benchmarks (GSM8K) than on less contaminated ones (FOLIO). Table 7 shows the signal is **mixed**: Qwen-2.5-1.5B and Phi-3 Mini show higher knockout faithfulness on GSM8K than FOLIO (+18.9pp and +25.3pp, consistent with contamination), while Llama-3.2 and Llama-3.1 show the reverse (−16.1pp and −23.8pp). This mixed pattern suggests contamination contributes to but does not solely explain the observed faithfulness patterns.

Per-difficulty MATH analysis. Within the MATH benchmark, knockout faithfulness *increases* with problem difficulty: Level 1 (easiest) shows 34% faithfulness while Level 5 (hardest) shows 42%. This is consistent with the contamination interpretation: easier problems are more likely memorized, leading models to bypass CoT perturbations and produce stored answers.

Probe cross-validation. Knockout and corruption probes agree on 68–89% of problems (Cohen's κ=0.29–0.39, fair agreement), confirming they measure a related underlying property rather than independent artifacts.

Within-family Qwen comparison. The controlled within-family analysis provides the strongest evidence because it eliminates architectural confounds. Table 8 shows that all six benchmark×probe deltas are negative: scaling from 1.5B to 7B consistently reduces faithfulness, with knockout drops of −15.4pp to −30.6pp and corruption drops of −13.4pp to −24.9pp. The absence of any reversal across benchmarks strengthens the tradeoff interpretation.

4.6 Paraphrase Stability

Figure 4 shows the distribution of CoT consistency scores across models. Paraphrase consistency

	GSM8K	MATH	FOLIO
KO Δ (7B−1.5B)	−28.0	−30.6	−15.4
COR Δ (7B−1.5B)	−13.4	−16.0	−24.9

Table 8: Within-family Qwen-2.5 comparison: faithfulness change (pp) from 1.5B to 7B. All deltas are negative across both probes and all benchmarks.

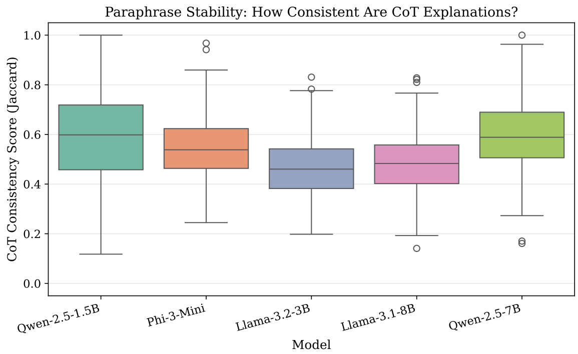

Figure 4: Paraphrase stability: distribution of CoT consistency scores across models. Medians range 0.47–0.60 with substantial within-model variance.

ranges from 47.2% (Llama-3.2) to 59.8% (Qwen-2.5-7B), with wide interquartile ranges (0.35–0.75) indicating high variability. Notably, there is *no meaningful tier effect* (small: 53.5%, medium: 54.1%), in contrast to the knockout/corruption probes. This suggests that CoT structural consistency is more a property of the task and prompt interaction than of model scale. Models produce slightly more consistent paraphrases for problems they solve correctly (+3–9pp).

4.7 Error Rates and Data Quality

JSON parse failures are a systematic issue for small models on hard benchmarks. Table 9 reports baseline error rates, which range from 0.5% (Qwen-2.5-7B and Llama-3.1 on GSM8K) to 42.0% (Qwen-2.5-1.5B on FOLIO). Error rows are excluded from all analyses; effective sample sizes range from 116 (Qwen-2.5-1.5B, FOLIO) to 399 (Qwen-2.5-7B, GSM8K). These exclusions could bias results if error-prone problems systematically differ in faithfulness; we note this as a limitation.

5 Discussion

A faithfulness–capability tradeoff. Within the 1.5–8B range, we observe that more capable models show lower perturbation faithfulness, suggesting that scaling improves answer re-

Model	GSM8K	MATH	FOLIO
Q-2.5-1.5B	1.5	26.5	42.0
Llama-3.2	8.0	16.5	5.5
Phi-3 Mini	9.0	41.5	25.0
Llama-3.1	0.5	17.5	15.0
Q-2.5-7B	0.5	9.5	5.5

Table 9: Baseline JSON parse error rates (%). Smaller models on harder benchmarks show substantially higher failure rates, reducing effective sample sizes.

trieval through mechanisms partly independent of CoT reasoning. The Qwen-2.5 within-family comparison—controlling for architecture—shows a 5$\times$ parameter increase yielding +30pp accuracy but $-$26pp knockout faithfulness, with all six benchmark$\times$probe deltas negative (Table 8). The per-benchmark analysis reveals this effect is concentrated on GSM8K (Cohen's h=-0.54) and less pronounced on FOLIO (h=0.002), suggesting that contamination is a contributing factor. However, the mixed contamination signal across model families (Table 7) and the universal within-family pattern argue that contamination alone is insufficient to explain the tradeoff. We emphasize this is observed within a narrow parameter range with n=5 models, and should not be read as a universal scaling law.

The anchoring effect. The gap between question-only (68–75%) and unbiased counterfactual (3–17%) answer-change rates reveals that CoT exerts a strong anchoring effect: models presented with stale reasoning largely ignore question modifications they would otherwise detect. CoT can thus *undermine* faithfulness by anchoring models to potentially outdated reasoning—a finding with practical implications for retrieval-augmented generation and multi-turn dialogue settings where context evolves. Practitioners should avoid feeding cached CoTs to models across context changes and should design evaluation pipelines that test for anchoring.

Implications for faithfulness evaluation. The 42–67pp prompt sensitivity in counterfactual probes (Cohen's h up to 1.85; Table 6) implies that faithfulness scores are substantially determined by evaluation methodology, not just model properties. We recommend future faithfulness evaluations include: (1) prompt sensitivity ablations, (2) no-CoT control conditions, (3) stratification by baseline correctness, and (4) per-benchmark reporting, since aggregate scores can mask benchmark-specific patterns.

Competent faithfulness vs. unfaithful bypassing. The conditional faithfulness asymmetry (Table 5) helps distinguish these alternatives. If models were primarily *reconstructing* removed steps from context (competent faithfulness), we would expect similar knockout rates regardless of baseline correctness. Instead, faithfulness is systematically *lower* on correctly-solved items (6–28%) than on incorrectly-solved items (18–54%), with the gap ranging from -11.6pp to -29.4pp across models. This suggests that for correctly-solved items, models largely bypass the provided CoT to retrieve memorized or prior-derived answers.

6 Limitations

- **Model scope**: Five models at Q4_K_M quantization ($\leq$9B), with only two medium-tier models. Mistral-7B was excluded due to high JSON parse error rates (baseline only completed); Gemma-2-9B was excluded due to infrastructure constraints.
- **Knockout validity**: Conflates reconstruction with unfaithfulness; the conditional faithfulness analysis (§4.3) partially mitigates but cannot fully resolve this without human annotation of model outputs.
- **Counterfactual prompt bias**: The consistency-checking instruction inflates scores (§4.4); the unbiased ablation quantifies but does not eliminate this confound.
- **Paraphrase metric**: Jaccard keyword similarity is a coarse structural proxy. The model generates its own paraphrases, risking circularity.
- **Reproducibility**: T=0.0 with seed 42 throughout; multi-seed robustness was infeasible. GPU non-determinism may cause minor variation.
- **Domain**: Mathematical and logical reasoning only; findings may not transfer to commonsense or factual reasoning.
- **Error exclusion bias**: JSON parse failures range from 0.5% to 42.0% (Table 9). If error-prone problems systematically differ in faithfulness, exclusion could bias results.

Ethical Considerations

Demonstrating CoT unfaithfulness has trust implications for users relying on model explanations. Our findings are specific to the five models studied and should not be overgeneralized. All experiments use open-source models locally; no human subjects or personal data are involved.

References

Wenhu Chen, Xueguang Ma, Xinyi Wang, and William W Cohen. 2023. Program of thoughts prompting: Disentangling computation from reasoning for numerical reasoning tasks. *Transactions on Machine Learning Research.*

Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, Christopher Hesse, and John Schulman. 2021. Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168.*

Luyu Gao, Aman Madaan, Shuyan Zhou, Uri Alon, Pengfei Liu, Yiming Yang, Jamie Callan, and Graham Neubig. 2023. PAL: Program-aided language models. *International Conference on Machine Learning.*

Shahriar Golchin and Mihai Surdeanu. 2024. Time travel in LLMs: Tracing data contamination in large language models. *Proceedings of ICLR.*

Simeng Han, Hailey Schoelkopf, Yilun Zhao, Zhenting Qi, Martin Riddell, Luke Benson, Lucy Sun, Ekaterina Zubova, Yujie Qiao, Matthew Bishi, and 1 others. 2022. FOLIO: Natural language reasoning with first-order logic. In *arXiv preprint arXiv:2209.00840.*

Dan Hendrycks, Collin Burns, Saurav Kadavath, Akul Arora, Steven Basart, Eric Tang, Dawn Song, and Jacob Steinhardt. 2021. MATH: Measuring mathematical problem solving with the MATH dataset. *NeurIPS Datasets and Benchmarks.*

Alon Jacovi and Yoav Goldberg. 2020. Towards faithfully interpretable NLP systems: How should we define and evaluate faithfulness? In *Proceedings of ACL*, pages 4198–4205.

Tamera Lanham, Anna Chen, Ansh Radhakrishnan, Benoit Steiner, Carson Denison, Danny Hernandez, Dustin Li, Esin Durmus, Evan Hubinger, Jackson Kernion, and 1 others. 2023. Measuring faithfulness in chain-of-thought reasoning. *arXiv preprint arXiv:2307.13702.*

Qing Lyu, Shreya Havaldar, Adam Stein, Li Zhang, Delip Rao, Eric Wong, Marianna Apidianaki, and Chris Callison-Burch. 2023. Faithful chain-of-thought reasoning. In *Proceedings of IJCNLP-AACL.*

Letitia Parcalabescu and Anette Frank. 2024. On measuring faithfulness or self-consistency of natural language explanations. In *Proceedings of ACL.*

Ansh Radhakrishnan, Karina Nguyen, Anna Chen, Carol Chen, Carson Denison, Danny Hernandez, Esin Durmus, Evan Hubinger, Jackson Kernion, Kamile Lukosuite, and 1 others. 2023. Question decomposition improves the faithfulness of model-generated reasoning. *arXiv preprint arXiv:2307.11768.*

Melanie Sclar, Yejin Choi, Yulia Tsvetkov, and Alane Suhr. 2024. Quantifying language models' sensitivity to spurious features in prompt design. In *Proceedings of NAACL.*

Miles Turpin, Julian Michael, Ethan Perez, and Samuel Bowman. 2024. Language models don't always say what they think: Unfaithful explanations in chain-of-thought prompting. *Advances in Neural Information Processing Systems*, 36.

Boshi Wang, Sewon Min, Xiang Deng, Jiaming Shen, You Wu, Luke Zettlemoyer, and Huan Sun. 2023a. Towards understanding chain-of-thought prompting: An empirical study of what matters. *Proceedings of ACL.*

Xiaohan Wang and 1 others. 2024. A survey on knowledge distillation of large language models. *arXiv preprint arXiv:2402.13116.*

Xuezhi Wang, Jason Wei, Dale Schuurmans, Quoc Le, Ed Chi, Sharan Narang, Aakanksha Chowdhery, and Denny Zhou. 2023b. Self-consistency improves chain of thought reasoning in language models. *arXiv preprint arXiv:2203.11171.*

Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Brian Ichter, Fei Xia, Ed Chi, Quoc Le, and Denny Zhou. 2022. Chain-of-thought prompting elicits reasoning in large language models. *Advances in Neural Information Processing Systems*, 35:24824–24837.

Xi Ye and Greg Durrett. 2022. The unreliability of explanations in few-shot prompting for textual reasoning. In *Advances in Neural Information Processing Systems*, volume 35.